

```
In [1]: #Importing required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import json as r
```

```
In [2]: #importing important libraries for prediction
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

#loading the dataset file =

```
r.load("http://localhost:8888/edit/Diabetes/diabetes_dataset.csv"
(http://localhost:8888/edit/Diabetes/diabetes_dataset.csv")) data = pd.read_csv(file)
```

```
In [3]: #Loading the dataset
data = pd.read_csv("diabetes_dataset.csv")
```

```
In [4]: data
```

Out[4]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedi
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

768 rows × 9 columns



```
In [12]: data.shape
```

Out[12]: (768, 9)

```
In [13]: data.columns
```


```
Out[13]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',  
              'Insulin',  
              'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],  
             dtype='object')
```

```
In [14]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 768 entries, 0 to 767  
Data columns (total 9 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   Pregnancies                          768 non-null    int64  
1   Glucose                              768 non-null    int64  
2   BloodPressure                        768 non-null    int64  
3   SkinThickness                        768 non-null    int64  
4   Insulin                              768 non-null    int64  
5   BMI                                  768 non-null    float64  
6   DiabetesPedigreeFunction             768 non-null    float64  
7   Age                                  768 non-null    int64  
8   Outcome                              768 non-null    int64  
dtypes: float64(2), int64(7)  
memory usage: 54.1 KB
```

```
In [5]: #checking for missing values  
sns.heatmap(data.isnull())
```

Out[5]: <Axes: >




```
In [6]: #Co relation matrix
correlation = data.corr()
print(correlation)
```

	Pregnancies	Glucose	BloodPressure	Skin
Thickness \				
Pregnancies	1.000000	0.129459	0.141282	
-0.081672				
Glucose	0.129459	1.000000	0.152590	
0.057328				
BloodPressure	0.141282	0.152590	1.000000	
0.207371				
SkinThickness	-0.081672	0.057328	0.207371	
1.000000				
Insulin	-0.073535	0.331357	0.088933	
0.436783				
BMI	0.017683	0.221071	0.281805	
0.392573				
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	
0.183928				
Age	0.544341	0.263514	0.239528	
-0.113970				
Outcome	0.221898	0.466581	0.065068	
0.074752				

	Insulin	BMI	DiabetesPedigreeFunction
on \			
Pregnancies	-0.073535	0.017683	-0.0335
23			
Glucose	0.331357	0.221071	0.1373
37			
BloodPressure	0.088933	0.281805	0.0412
65			
SkinThickness	0.436783	0.392573	0.1839
28			
Insulin	1.000000	0.197859	0.1850
71			
BMI	0.197859	1.000000	0.1406
47			
DiabetesPedigreeFunction	0.185071	0.140647	1.0000
00			
Age	-0.042163	0.036242	0.0335
61			
Outcome	0.130548	0.292695	0.1738
44			

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
In [7]: ▶ plt.figure(figsize=(12, 10))
p=sns.heatmap(correlation, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



```
In [8]: ▶ X =data.drop("Outcome",axis=1)
Y =data['Outcome']
X_train,X_test,Y_train,Y_test =train_test_split(X,Y,test_size=0.2)
```

```
In [9]: ▶ model = LogisticRegression(max_iter = 1000) #Adjusting the number of iterations
model.fit(X_train, Y_train)
```

Out[9]:

LogisticRegression (https://scikit-learn.org/1.4/modules/generated/sklearn.linear

LogisticRegression(max\_iter=1000)

```
In [10]: ▶ #Making Prediction
          predictions = model.predict(X_test)
          print(predictions)
```

```
[1 0 0 0 0 1 0 0 0 0 0 0 1 0 0 1 0 0 0 0 1 0 1 1 0 1 0 1 0 0 0 0 0
0 1 0
 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1
0 0 0
 1 1 0 0 0 1 0 1 0 1 0 0 0 1 1 1 0 0 1 0 0 1 0 0 0 0 0 0 0 1 1 0 0 1
0 0 0
 1 0 1 1 0 1 1 0 0 0 0 0 1 0 1 0 0 1 0 1 1 1 1 0 0 1 0 0 0 0 0 0 1 0
0 1 1
 0 1 0 0 0 1]
```

```
In [15]: ▶ accuracy = accuracy_score(predictions,Y_test)
          print(accuracy)
```

```
0.7987012987012987
```

**This (accuracy) means that the model correctly predicted the outcome of diabetes in approximately 79.87% of the cases in the test set**